

A photo you can't find is a memory you've lost

1.5B+

people use Google Photos each month

9T+

photos and videos stored

370M+

people search their photos each month (1 in 4)

150M

Google One subscribers paying for storage

More people finding the photo keeps the library worth paying for

Our goal: more users retrieve a photo they remember but can't precisely describe. **Why Google cares:** a library people can't search is worth less to keep, and storage is what Google One sells.



It becomes one number: User Retrieval Rate (URR)

Of the people who search for a photo they can only vaguely describe, the share who reach it within 28 days. It counts people, not searches.

$$URR = E \times [1 - (1 - I \times S \times R)^n]$$

E Express it

Can they put the memory into words at all?

Today: a search box, or a question to Ask Photos

I Be understood

Does Photos read the clue the way they meant it?

Today: matches people, places and things

S See it surface

Is the photo in the results at all?

Today: a grid of matching photos

R Recognise it

Can they tell it is the right one?

Today: thumbnails with date and place

n Try again

How many tries before they give up?

Today: retype the search

How to read it: a person succeeds only if they can express the memory and one try passes I, S and R together. One weak step caps the whole number, so the next slides find which step breaks.

We turned public complaints into a structured map of failed retrievals

1 · COLLECT

85,140

Public posts
(Play Store, App Store, YT, Reddit)

2 · SCREEN

819

Relevant (Gate A 5,305 → Gate B 1,333 sample)

3 · EXTRACT

144

Specific attempts (720 episodes; 62 scoreable)

4 · AUDIT

203

Pairs checked blind by Qwen 27B model

5 · COMPARE

9 Areas

Ranked only on agreed fields (O1–O9)

Extracted Field in Schema	The Brief's Underlying Question Answered	Second-model agreement
asset_type	What kinds of old photos do users struggle to retrieve?	Validated in Gate B
cues_retained (14 cue types)	What do people actually remember about the photo?	Jaccard 0.557
cues_lost (date, place, album, words)	What have they forgotten when search fails?	Jaccard 0.673
query_verbatim · search_mode	How do they search when memory is incomplete?	Quote verify 85.2%
failure_stage (5 stages)	Where does retrieval break? (Part 3 decomposition)	κ 0.509 (Moderate)

The audit, not the model, set the ranking

1. **Structure, not sentiment:** Every review was converted into a structured, typed record with fixed vocabularies, not qualitative impression tags.

2. **Pre-registered rules:** Hypotheses H1–H5 and decision thresholds were set before extraction, so the data could not be read to fit them.

3. **Our audit changed the direction:** it did not reliably back our first hypothesis ranking (Jaccard 0.347), but it did reliably find where retrieval broke (κ 0.509). So we switched to failure stages.

retrieval-discovery-engine.vercel.app

People remember the episode, but Search needs the date

Engine and survey agree: the time is kept, the date is lost

Question	Engine (144 Specific Attempts)	Survey (n=15 / 14 Dedup)
Photo Kinds	Personal photos 80 · multi-photo 26 · videos 16 · screenshots 7 · docs 3	All 3 real-trouble cases were document or medicine photos
Remembered	Roughly when 40 · object 17 · exact date 14 · text 12 · people 10	Roughly when 8 · object 7 · people with me 6
Forgotten	The date 37 · album 29 · exact words 10 · place 10	When it was taken 9 · words to search 9
Search Method	20 of 32 queries are 1 word ("dog"); 10 name a time ("Halloween 2024")	First move: scroll 9/15; own-app: 4 of 5 strip time to bare nouns

O1 is the largest group with a memory search ignores

O1 (Rough time or an event): 42 → 37 after exclusions (18 never surfaced, 12 misread).

O8 (No clue at all): 47 attempts, larger, but there is no memory for search to use. O1 is the largest group where memory exists and search ignores it.

Excluded: O7 exact date known (14) · O9 path moved by an update (12).

Episodic time wins; dead-end recovery loses

H1 Episodic time: **SUPPORTED (root cause)**.

Roughly when is the #1 kept cue (40); the date is the #1 lost cue (37).

H2 Recognition: **SUPPORTED (secondary)**. 41.7% never surfaced; 7 of 15 ended unsure.

H3 Dead-end recovery: **REFINED, NOT THE LEAD**.

Overruled by the audit: recovery is 0.7% of failures (stage κ 0.509).

H4 Hinglish: **NOT TESTED**. Weak signal (2 of 15 survey, 1 of 6 tests); kept as a design choice.

H5 Text not indexed: **WEAKLY SUPPORTED** (3 docs; the index scores 0.000 on text in images).

H6 Path changed: **PRESENT BUT MINOR** (8.3%), added post hoc.

77% fail before any retry: the clue is misread or the photo never appears

S Never surfaced	60 (41.7%)
I Clue misread	51 (35.4%)
E · R · retry	2 · 2 · 1

When they can't find it, people scroll—and nearly half still leave unsure

A directional survey: 15 stories we could score
Structured self-serve questionnaire (n=15, 23–30 Sep): every answer maps onto the engine's vocabulary.
Why: 100% scoreable stories (vs 8.6% of public posts); reached 15 people vs 2 call contacts.
Costs: self-reported convenience sample; it shows patterns, not how common they are.

[Survey questions ↗](#)

The survey repeats the engine's pattern

Metric	Engine (144)	Survey (15 / 14)
Most remembered	Roughly when (40)	Roughly when (8)
Ended unsure	Not measurable	7 of 15 (47%)
First move: scroll	15 workarounds	9 of 15 (60%)

It costs most people over five minutes
Trips and outings, a photo for LinkedIn, a medical report. **12 of 15 spent 5+ minutes**, past the brief's five-minute bar; 3 had to ask someone or get the document again.

4 of 5 kept roughly when; only 1 put it in the search

ID	Target Photo	What They Remembered	What They Typed	Break	Outcome
S04	Wedding photo	Roughly when, event, people	"wedding, pichle saal diwali"	Interpretation	? Unsure
S03	Personal photo	Roughly when, event, object	"gym"	Recognition	? Unsure
S08	Personal photo	Roughly when, event	(Didn't know what to type; scrolled)	Expression	✓ Found
S12	Personal photo	Roughly when, event	(Never typed; opened album)	Expression	? Unsure
S06	Medicine bill	Object, text in bill	"medicine, bill"	Surfacing	✗ Never found

"wedding, pichle saal diwali" — S04 (Hinglish: a festival and relative year, not a calendar date)
"medicine, bill" — S06 (Never found; had to get the document again)
Ask Photos: 7 of 15 had never heard of it; all 4 who used it got "related photos, but not the one".

Our target user remembers "when," but not the calendar date

Defined by what they remember, not by age
Long-time Google Photos users searching for a personal moment they can place only roughly.
They remember "around Diwali", "my sister's graduation", "last winter".

Segment Sizing Filter	Engine (144)	Survey (15)
Kept rough time or event (O1)	42 (29%)	8 (53%)
Target Segment (After Exclusions)	37 (26%)	8 (53%)
Personal photo or video	30 of 37	All 8
Large library (5,000+ items)	Not measured	5 of 8

42 → 37: minus 2 who also knew the exact date (plain search works) and 3 whose browse path moved in an update (H6).

Why not an age band? Our evidence points to memory state and library habits, not age, as the driver. We will use age as a diagnostic cut in the A/B, not as the targeting rule.

Memory stores episodes; Photos stores items and dates



"I'll type in something super simple, like 'Halloween 2024' and it seriously can't find anything?" — Play Store review

- Asking why, for the 37 segment attempts:
- 1. Why isn't the photo found? It never surfaces (18 of 37).
 - 2. Why not? Search misreads the clue (12): "last Diwali" is neither a keyword nor a date.
 - 3. Why does that matter? People keep the event and lose the date (the #1 lost cue, 37 of 144).
 - 4. Root cause: nothing turns "roughly when" into a date window.
Real app: 4 of 5 testers typed bare nouns ("vacation", "gym"); 0 of 5 found the photo.

We recommend helping users recognise the moment—not write a better query

The problem, in the user's words

"I remember the event and roughly when ("last Diwali"), but Photos wants a date or keyword I've forgotten. So I scroll by hand, and even then I'm not sure the photo I found is the one."

We read this as a memory-to-recognition problem, not a generic search-quality problem.

The workaround is the design blueprint

Engine: 15 of 21 workarounds are manual scrolling (7 of 10 failed).

Survey: scrolling was the first move for 9 of 15; 5 of those 9 still ended unsure.

We recommend Photos take this step for them: turn the rough time into a date window and group the moments around it.

Our decision: solve the memory-to-recognition gap, not redesign Search for every query.

Each option helps, but none is built for "I'll know it when I see it"

Option	Good at	Gap for a vague memory
Photos search	People, places, things, dates	Memory must become search terms
Ask Photos	Plain-language questions	All 7 who tried: "related, not the one"
Apple Photos	Describing the photo	Describes items, not when it was
Memory Trails	Rough time → moments	Recognise, then recover the moment

Ask Photos helps people ask; we help them spot the right moment when its answer is only related.

Hypothesis Why Google can win: years of personal-library signals plus a recognition-first interaction, not the data alone. Rivals could build it too, so we must validate it.


Five times the evidence changed the plan

1. **Metric:** per-search success → **per-user URR**
2. **Outcomes:** recovery agent → **5 stages + browse**
3. **Discovery:** H3 ranked top → **overturned by audit**
4. **Reported behaviour:** retry search → **scroll by hand**
5. **Problem:** "search is bad" → **items vs. moments**

Memory Trails turns a vague memory into a trail of recognisable moments

We recommend placing it inside Search, where it fails

It doesn't ask for a perfect description; it helps you recognise the right moment from a short trail. A failed search offers "Can't describe it?" with the words carried over.



memory-trails-v2.vercel.app

Blue frame = AI step

Our MVP: one AI step per failure; the user confirms

Only a rough time → turns it into a date window

Mixed clues → editable chips, plus "I also heard: + dosa"

A flat grid → moments; full matches rank first

Why this one? → shows matched vs approximate clues

Wrong moment → nearby moments; undo a rule-out

It never confirms for you, and can't always find the photo.

Big gain with rich memory, modest with one vague cue Modelled

Memory kept	What the query has	Baseline photo top 20	Ours photo top 20	Ours moment top 5
L3 (Fully described)	Rough time, place, album word	0.172	0.962	0.967
L2 (Two vague cues)	Scene + 2 vague clues	0.209	0.276	0.333
L1 (One vague cue)	Scene + rough time	0.242	0.309	0.333
L0 (Content only)	Scene only	0.312	0.312	0.267
Weighted by real memory	Mixed as in the 144 real attempts	0.259	0.334 (+0.075)	0.338 (+0.113)

Read this first: L3's 0.962 needs three clues; only 6 of 144 real attempts kept that many. Weighted by real memory: photo **+0.075**, moment **+0.113**. Grouping helps more than ranking.

Scope: the interaction on 1,282 CC test photos, not production accuracy. Not a Search redesign; no unasked sensitive memories or stored life-event profile.

The MVP helped users find moments—but compound time still breaks it

ID	Own Photos search → result	Task 1 (cat)	Task 2 (dog)	Sure	vs GP	Would use
R01	"Wedding before COVID" → wrong years	✓ Found, sure	✗ Not found	5/5	5/5	Every time
R02	"gym" → can't tell which	✓ Found, sure	✓ Found (1-tap alt)	4/5	4/5	When fails
R03	(Apple Photos user)	✓ Found, sure	✓ Found (nearby)	5/5	—	Every time
R04	"vacation" → wrong years	✓ Found, sure	✓ Found (1-tap alt)	4/5	4/5	When fails
R05	"Wedding event" → can't tell which	✓ Found, sure	✓ Found (clue edit)	4/5	4/5	Every time
R06	"Kerala photos" → can't tell which	? Right day, unsure	✓ Found (1-tap alt)	4/5	3/5	Every time

6 of 6

Fit the segment

4.33 / 5

Task 1 sureness (self-reported)

4.0 / 5

Ease vs Photos

0 of 5

Found it in own Photos

MVP test questions ↗

memory-trails-v2.vercel.app ↗

Testers valued the moment, not a grid

- "Shows the moment, not a grid" — R02
- "Moving to nearby months, like scrolling but faster" — R03
- "Works without knowing what to type" — R04
- "Finds exact month & year... understands mix of Hindi and English" — R06

From these six tests, we will fix three things next

- Compound time:** R01 stuck on "pichle ke pichle saal" (year before last). *Fix:* parse compound offsets; editable year.
- Clues missed:** R04 didn't notice them. *Fix:* high-contrast "Showing Diwali 2025 · Tap to change".
- Speed:** R05, R06 wanted faster loading. *Fix:* pre-cache moments.

We will test whether Memory Trails increases confirmed retrieval

One number: users who reach the photo in 28 days

Proposed

Definition: share of users with a vague-memory search in 28 days who reach the photo. Counts users, not searches.

Baseline 40% → target 45%. The baseline is a self-reported proxy (6 of 15 sure they found it). +5 pts is the smallest lift worth shipping.

We propose an 8-week user-level A/B, not yet run: ~1,600 users per arm for 80% power at α 0.05.

What the MVP has shown so far (not the A/B)

Modelled right photo in top 20: 0.259 → 0.334; right moment in top 5: +0.113

Proxy 6 testers' "sure" score 4.33 / 5: self-reported, not production URR

We decide now when to ship, extend or stop

Proposed

Ship if URR rises ≥5 pts and no guardrail trips. Extend if the lift is positive but under 5. Stop if false confirmation passes 3%.

Three numbers decide it

Proposed

Role	Metric	Target
Primary	Confirmed retrieval (URR)	40% → 45%
Guardrail	False confirmation: said "found it", then kept searching	<3%
Experience	Time to first moment (95th percentile)	<1.5 s

Secondary signals, to explain any movement

Entry after a failed search (≥12%) · clue edits (15–25%) · top card matches all clues (rising) · confirm after opening a moment (≥55%) · "Not in any of these?" (falling) · abandonment (no rise)

The A/B has not run yet. Our evidence supports the direction but not a production-scale lift. We will test whether Memory Trails lifts confirmed retrieval without raising false confidence.

The opportunity is real, but we need one controlled test before scaling

Each risk has a mitigation built or planned

Risk	Mitigation & evidence
R1. The win sits where memory is richest (+0.79 at L3, +0.07 at L1)	Lean on moment recognition (right moment in top 5: 0.333 at L1–L2); report L1 separately in the A/B.
R2. Date filters hide photos	Soft scoring built: a wrong date lowers rank, never hides (L3: 0.851 → 0.962).
R3. False confirmation (worst failure)	Never auto-confirm; show the surrounding photos; stop launch above 3%.
R4. Can't read text in photos (0.000 on receipts)	Add OCR text to the index.
R5. Ask Photos may already suffice	4 of 4 survey and 3 of 3 test users got "related photos, not the one".
R6. Sensitive memories (medical, relationships)	User-initiated only; the session persists nothing.

The evidence is small and mostly self-reported

- Play Store is 86.2% of engine posts
- Survey n = 15 vs a target of 30, self-reported
- MVP tests n = 6, unmoderated, no session logs
- URR baselines modelled; no telemetry
- Eval library: 1,282 CC photos, synthetic episodes
- Ranking weights are judgement
- Hinglish (H4) is a weak signal only

Our next steps, and what we ask for

1. Parse compound dates ("pichle ke pichle saal")
 2. Make the active date clue unmissable
 3. Pre-cache moments so they load faster
- Started: a native app build (pinch-zoom, swipe-down, haptics), not yet user-tested.

We ask for approval to fix these three observed issues and run a one-market A/B test.